

Finding Change in the CouchDB

Dr. James E. Marca

January 13, 2013

Big Data



We've always had big data!

- ☐ We've been handling big data for years in transportation
- ☐ What's wrong with sticking to the old tools?

Old Way:

- ☐ PostgreSQL
- ☐ Simple queries
- ☐ Canned aggregates

“Is the current reading of volume and occupancy from some detector more or less than what should be expected?”

But we had a new server...

- ☐ Lots of CPU
- ☐ Big Disks
- ☐ Raw data already loaded
- ☐ in PostgreSQL...



- ☐ The index was too big
- ☐ Extensive disk swapping
- ☐ Super-duper slow!

Origins of Big Data and NoSQL

- ☐ Google (Big Table)
- ☐ Amazon (Dynamo)





Some options available now

- ☐ Hadoop
- ☐ MongoDB
- ☐ Riak
- ☐ CouchDB
- ☐ BigCouch
- ☐ Couchbase
- ☐ Cassandra
- ☐ Voldemort
- ☐ RethinkDB
- ☐ Datalog
- ☐ and more all the time...



CouchDB is great for transportation data



Because transportation data is


- ☐ largely observations and measurements
- ☐ write once
- ☐ read raw data for short term applications
- ☐ process raw data into summary stats



Traffic events



What is CouchDB?



CouchDB is a document oriented database

- ☐ no schema
- ☐ JSON documents
- ☐ Views (map and reduce steps) are analogous to queries
- ☐ Eventual consistency

The CAP Theorem



The CAP theorem describes a fundamental limitation of databases

- ☐ **Consistency** All database clients see the same data, even with concurrent updates.
- ☐ **Availability** All database clients are able to access some version of the data.
- ☐ **Partition Tolerance** The database can be split over multiple servers.

Choose any two

Traditional RDBMS

- ☐ Consistency
- ☐ Availability

CouchDB



Availability



Partition Tolerance

“Eventually Consistent”

Consistency isn't that big a deal for traffic data


- This isn't stock trading or e-commerce
 - It's okay if:
 - you are 30s behind reading a loop
 - Controller A has slightly more current info than Controller B
- A universally consistent view isn't mission critical

Bonus: CouchDB has master-master replication

Replication is super awesome!

20

Replication is the *change* I was looking for



We can have distributed databases!!

- ☐ Put the database *at the detector*
- ☐ Move data around by using replication
- ☐ A TMC can pull from *all* detector databases
- ☐ A Local TMC can *limit replication* to relevant detectors
- ☐ A traveler can replicate traffic DBs along common routes

Relax



23

Practical experience

1. Processing, storing raw loop detector data
2. Imputing missing detector data
3. A single stash for storing detector metadata

1: Processing Raw Loop Detector Data

- ❑ Orange County, California (CalTrans District 12)
- ❑ about 900 mainline detectors
- ❑ Process in R
 - compute 27 different measures per location
 - for 20 minute running time window
 - (vol, occ per lane + 27) per 30s period
 - run models estimating relative risk of accident types
- ❑ 280GB of data (three years)
- ❑ 590GB of generated views

2: Storing the results of imputation runs

- ❑ 400GB of data
- ❑ 700GB of generated views
- ❑ Databases spread over three machines
- ❑ Uses per-district collation databases
- ❑ Analysis step used CouchDB to coordinate multiple processes
 - Local “state” database on each machine
 - State databases replicated with each other
 - No overlapping runs were observed

3: Convenient stash for detector metadata

- ☐ 20GB of data
- ☐ 222MB of generated views
- ☐ Uses GeoCouch extensions, stores location of each detector
- ☐ Stashes all known metadata about each detector in a single place
- ☐ Uses binary attachments to save R analysis output (plots, files)
- ☐ Small enough to replicate to my laptop too

Questions?

James E. Marca

UCI ITS

jmarca@translab.its.uci.edu

Research funded by Caltrans, California Air Resources Board

Database design:

- ☐ One CouchDB database per detector
- ☐ One document per day

Document per day reasons:

- ☐ Based on ~~painful experience~~ informal testing
- ☐ One document per year is too big to process
- ☐ One document per timestamp would work okay,
- ☐ But the web *application* uses daily data
- ☐ CouchDB sorts by document id, id is based on timestamp
- ☐ HTTP GET:
/vdsdata/d12/2007/1202248/1202248 2007-01-03 00:00:00

30

Database per detector per year reasons:

- ☐ One big database is possible, but
- ☐ Impossible to *split* or *shard* over different machines (now can use BigCouch)
- ☐ makes better use of multi-core machines when generating views

Use views to run models and summarize data

- Views are CouchDB's version of map/reduce
- Write JavaScript code for the map function that is run on each document (to apply models, run summaries, etc)
- ProTip[©] Only use embedded Erlang reduce functions like `_count`, `_sum`, and `_stats`

Difficulty: Need to use another database to collate model output

- ☐ Pipe summaries of the per-detector views to a single database for all detectors in the district
- ☐ Requires external programming
- ☐ Difficult to automate